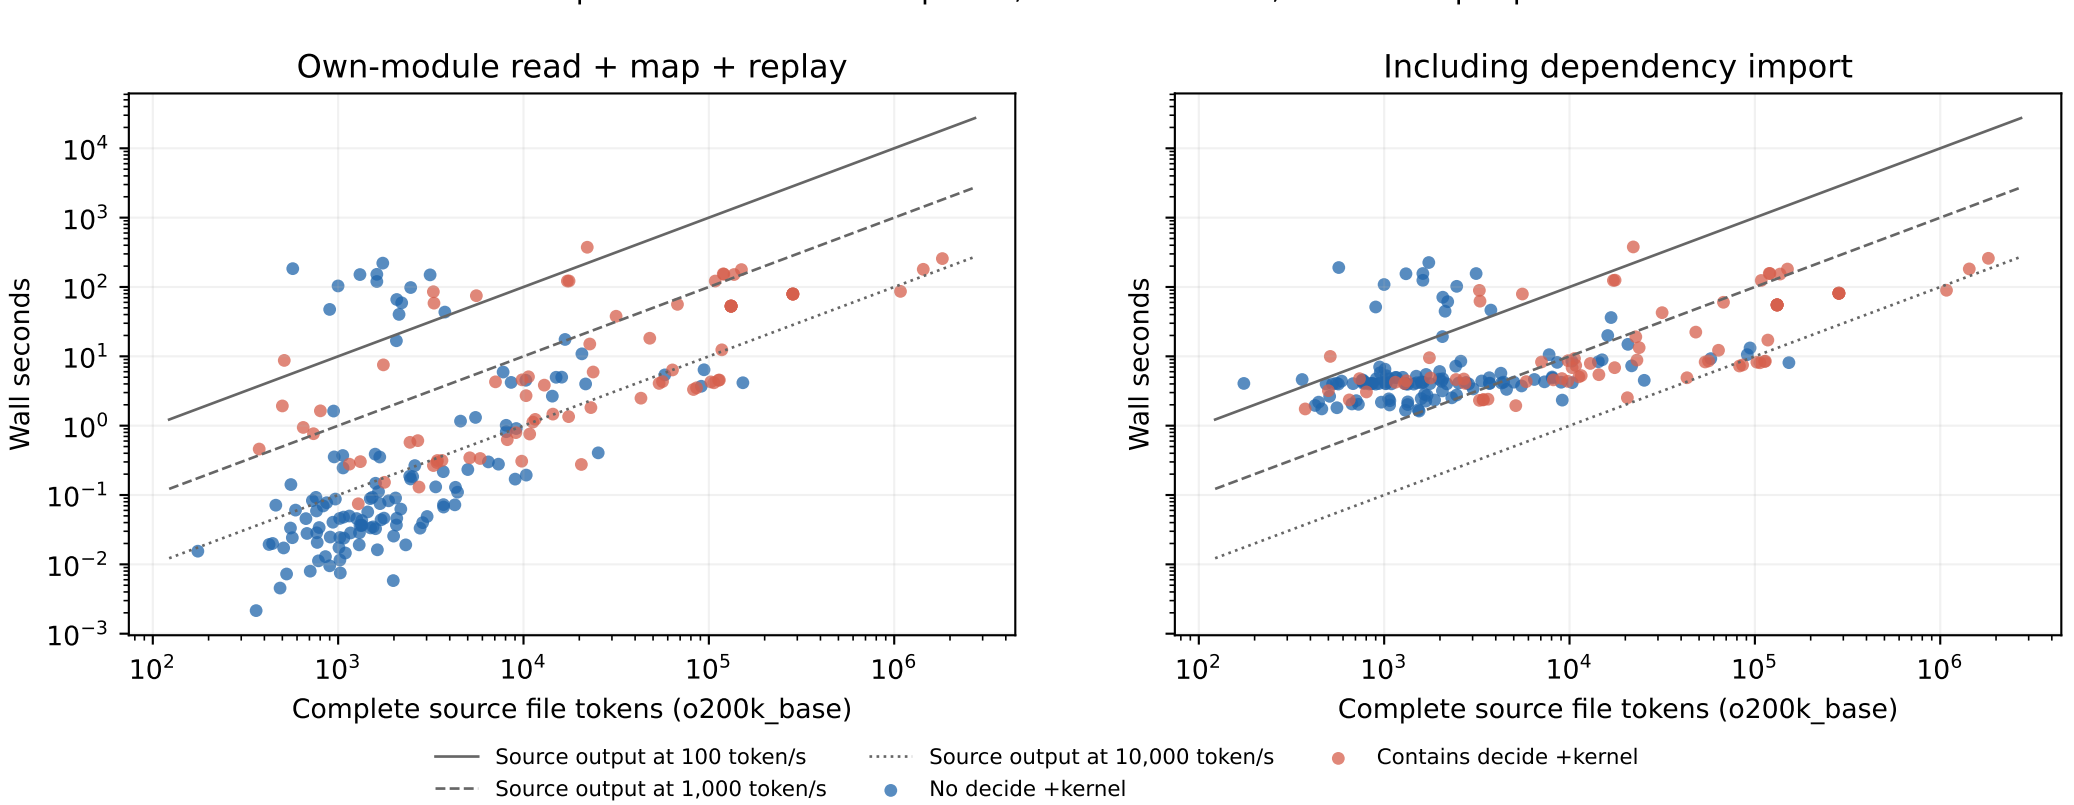
200 sampled modules: 200 completed, 0 time-censored; one P core per process
Own-module read + map + replay
Including dependency import
Wall seconds
Complete source file tokens (o200k_base)
Source output at 100 token/s
Source output at 1,000 token/s
Source output at 10,000 token/s
No decide +kernel
Contains decide +kernel
Fixed stratified sample, not population prevalence. Source-output lines are scenarios, not measured AI generation.